

CS5100: Foundations of Artificial Intelligence

Informed search algorithms

Dr. Rutu Mulkar-Mehta

Lecture 2

Administrative

- Assignment P0 – How many did it?
- Project 1- How many started it?

Review from Lecture 1

- Agents and Agent types
- Agents to perform search
- Uninformed Search Techniques
 - Breadth First Search
 - Depth First Search
 - Uniform Cost Search
 - Iterative Deepening Search

Terminology

- Start and Goal
- Parent Node
- Child Node
- Fringe or Frontier
 - Next list of nodes to be expanded
- Search Strategy
 - BFS, DFS, UCS etc.

Space Time Complexity

- Time: Number of nodes generated during the search
- Space: Maximum number of nodes stored in memory

Review: BFS

- Starts with the root node and explores all neighboring nodes
 - Repeats this for every one of those
- This is realized in a FIFO queue
- It does an exhaustive search until it finds a goal.

Review: BFS

- BFS is complete
 - (i.e. it finds the goal if one exists and the branching factor is finite).
- It is optimal (Finds best solution)
 - (i.e. if it finds the node, it will be the shallowest in the search tree).

Review: BFS

- Space: $O(b^d)$
- Time: $O(b^d)$
- b = branching factor
- d = depth to which BFS is reached

Review: DFS

- Explores one path to the deepest level and then backtracks until it finds a goal state.
- This is realized in a LIFO queue (i.e. stack).

Review: DFS

- DFS is complete
 - (if the search tree is finite).
- It is not optimal
 - (it stops at the first goal state it finds, no matter if there is another goal state that is shallower than that).

Review: DFS

- Space: $O(bm)$
 - Much lower than BFS
- Time: $O(b^m)$
 - (Higher than BFS if there is a solution on a level smaller than the maximum depth of the tree).
 - Danger of running out of memory or running indefinitely for infinite trees.
- b = branching factor
- m = maximum depth

Review: Iterative Deepening DFS

- The search depth for DFS increased iteratively over time, gradually increasing the maximum depth to which it is applied.
- This is repeated until finding a goal.
- Combines advantages of DFS and BFS.
- It is complete.
- It is optimal (the shallowest goal state will be found first, since the level is increased by one every iteration).

Review: Iterative Deepening DFS

- Space: $O(bd)$
 - (better than DFS, d is depth of shallowest goal state, instead of m , maximum depth of the whole tree).
- Time: $O(b^d)$
- b = branching factor
- m = maximum depth

Review: Uniform Cost Search

- Expand lowest path cost
- Demands the use of a priority queue

Uniform Cost Search

- Space: $O(b^{1+C/e})$
- Time: $O(b^{1+C/e})$
- Can be much greater than b^d
- C = cost of optimal solution
- e = cost of every action
- b = branching factor
- m = maximum depth

Today

- Informed Search
 - Problem Specific Knowledge is provided
 - This is beyond the definition of the problem itself
 - This information can be used to find solutions more efficiently than uninformed search



Today

- Best-first search
- Greedy best-first search
- A* search



Best-first search

- Implementation:
 - Order the nodes in fringe in decreasing order of desirability
- Idea: use an **evaluation function** $f(n)$ for each node
 - estimate of "desirability"
 - Expand most desirable unexpanded node
 - $f(n)$ is based on the heuristic $h(n)$
- Special cases:
 - greedy best-first search
 - A* search

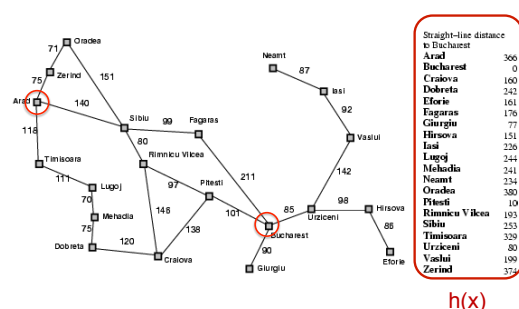
Search Heuristics

A heuristic is:

- A function that *estimates* how close a state is to a goal
- Designed for a particular search problem
- Examples: Manhattan distance, Euclidean distance for pathing



Romania with step costs in km




Greedy Search



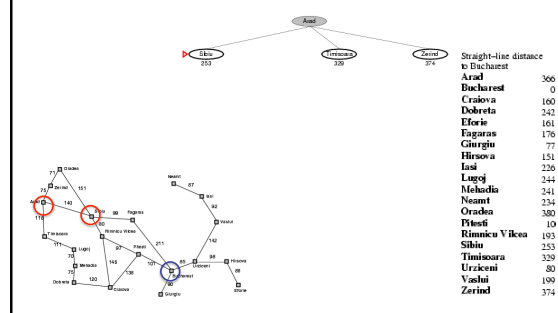
Greedy best-first search

- Evaluation function $f(n) = h(n)$ (heuristic)
 - e.g. estimate of cost from n to goal
- e.g., $h_{SLD}(n)$
 - straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal

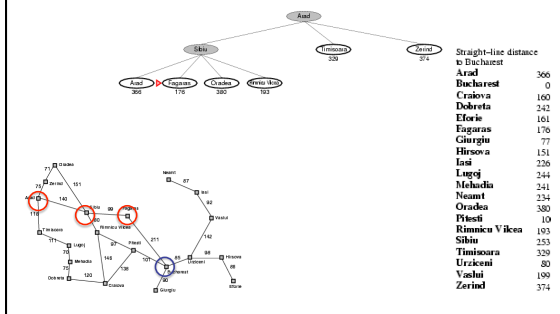
Greedy best-first search example



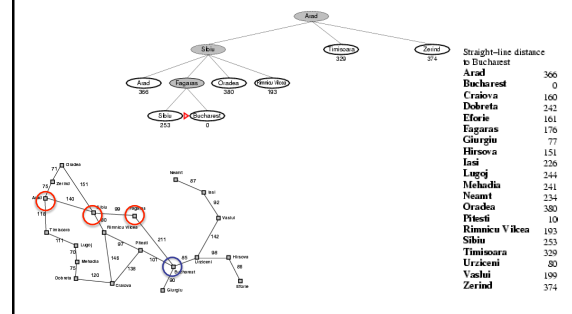
Greedy best-first search example



Greedy best-first search example




Greedy best-first search example



Properties of greedy best-first search

- **Complete?** No
 - can get stuck in loops,
 - e.g., Iasi → Neamt → Iasi → Neamt →
- **Time?** $O(b^m)$
 - but a good heuristic can give dramatic improvement
- **Space?** $O(b^m)$
 - keeps all nodes in memory
- **Optimal?** No

Best First: Incomplete



A* Search



A* search


- Idea: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to goal
 - $f(n)$ = estimated total cost of path through n to goal

A* search and UCS


- A* search is identical to UCS, except it uses $g(n) + h(n)$ instead of $h(n)$

Combining UCS and Greedy

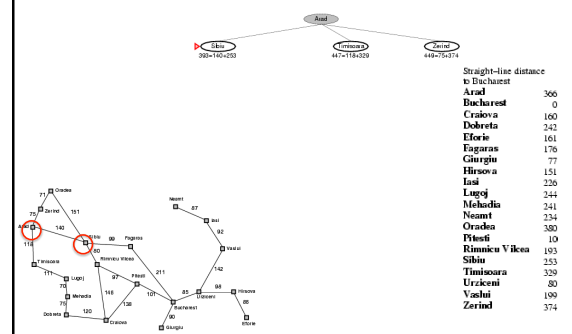
- Uniform-cost orders by path cost $g(n)$
- Greedy orders by goal proximity $h(n)$



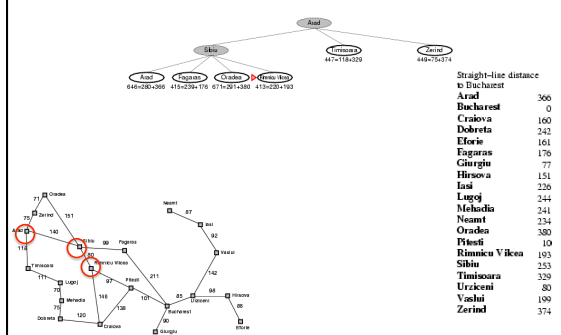
A* search example



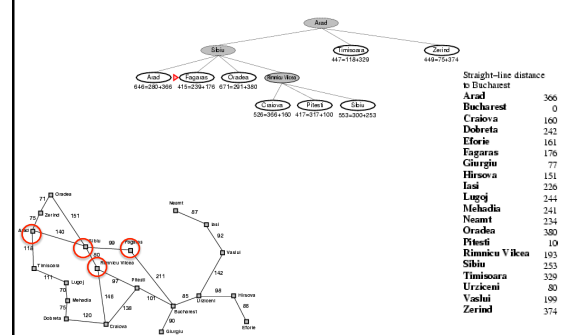
A* search example



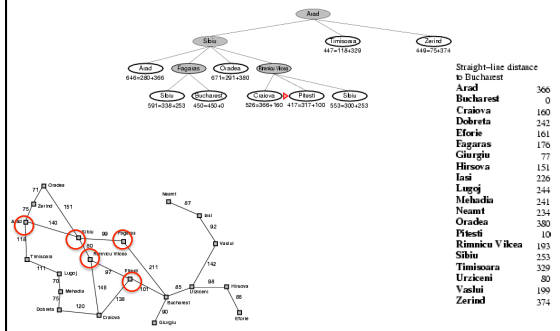
A* search example



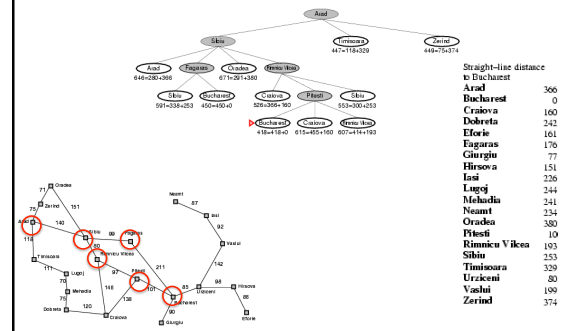
A* search example



A* search example




A* search example




When should A* terminate?

- Should we stop when we enqueue a goal?




- No: only stop when we dequeue a goal

Is A* Optimal?




- What went wrong?
- Actual bad goal cost < estimated good goal cost
- We need estimates to be less than actual costs!


Admissible Heuristics



Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs


Admissible Heuristics

- A heuristic h is **admissible** (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$


where $h^*(n)$ is the true cost to a nearest goal

- Examples:



- Coming up with admissible heuristics is most of what's involved in using A* in practice.

Optimality of A* Tree Search




Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:


- A will exit the fringe before B



Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B



- $f(n)$ is less or equal to $f(A)$


$$\begin{aligned} f(n) &= g(n) + h(n) && \text{Definition of f-cost} \\ f(n) &\leq g(A) && \text{Admissibility of } h \\ g(A) &= f(A) && h = 0 \text{ at a goal} \end{aligned}$$

Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B

- $f(n)$ is less or equal to $f(A)$
- $f(A)$ is less than $f(B)$




$$\begin{aligned} g(A) &< g(B) && \text{B is suboptimal} \\ f(A) &< f(B) && h = 0 \text{ at a goal} \end{aligned}$$

Optimality of A* Tree Search: Blocking

Proof:


- Imagine B is on the fringe
 - Some ancestor n of A is on the fringe, too (maybe A!)
 - Claim: n will be expanded before B
- $f(n)$ is less or equal to $f(A)$
 - $f(A)$ is less than $f(B)$
 - n expands before B
- All ancestors of A expand before B
 - A expands before B
 - A* search is optimal



$$f(n) \leq f(A) < f(B)$$


Properties of A*

Properties of A*




UCS vs A* Contours

- Uniform-cost expands equally in all “directions”



- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



Video of Demo Contours (Empty) -- UCS



Video of Demo Contours (Empty) -- Greedy



Video of Demo Contours (Empty) -- A*



Video of Demo Contours (Pacman Small Maze) – A*



Comparison




Greedy

Uniform Cost

A*


Creating Heuristics



Creating Admissible Heuristics


- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available

366




- Inadmissible heuristics are often useful too

Example: 8 Puzzle




- What are the states?
- How many successors from the start state?

Example: 8 Puzzle



- What is the branching factor?
- If the average solution ~22 steps, what is the max depth that you would like to hit in a search tree?

Example: 8 Puzzle



- How many states?
 - $9!/2$
 - half the states are impossible to solve

Relaxed problems



- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem

8 Puzzle - heuristics

- $h_1(n)$: The number of Misplaced Tiles
 - If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
- $h_2(n)$: sum of distances of the tiles from their goal position
 - If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution

8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = 8$
- This is a **relaxed-problem** heuristic




Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
TILES	13	39	227

Statistics from Andrew Moore

8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?



- Total **Manhattan** distance
- Why is it admissible?
- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$

Average nodes expanded when the optimal path has...			
	...4 steps	...8 steps	...12 steps
TILES	13	39	227
MANHATTAN	12	25	73

8 Puzzle III

- How about using the **actual cost** as a heuristic?
 - Would it be admissible?
 - Would we save on nodes expanded?
 - What's wrong with it?
- With A*: a trade-off between quality of estimate and work per node
 - As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

A*: Summary



A*: Summary

- A* uses both path costs and heuristics
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems

